DataFrame Methods	Types	Functors	Random Generators
append_column(2)	enum class drop_policy { }	struct AutoCorrVisitor { }	gen_bernoulli_dist( )
append_index(2)	enum class fill_policy { }	struct BetaVisitor { }	gen_binomial_dist( )
bucketize( )	enum class int_dist_policy {}	struct CorrVisitor { }	gen_cauchy_dist( )
bucketize_async( )	enum class io_format { }	struct CovVisitor { }	gen_chi_squared_dist( )
create_column( )	enum class join_policy { }	struct CumMaxVisitor { }	gen_exponential_dist( )
drop_missing( )	enum class nan_policy { }	struct CumMinVisitor { }	gen_extreme_value_dist( )
fill_missing( )	enum class random_policy { }	struct CumProdVisitor { }	gen_fisher_f_dist( )
gen_datetime_index( )	enum class return_policy { }	struct CumSumVisitor { }	gen_gamma_dist( )
gen_sequence_index( )	enum class shift_policy { }	struct DotProdVisitor { }	gen_geometric_dist( )
get_col_unique_values( )	enum class sort_state { }	struct GroupbySum { }	gen_lognormal_dist( )
get_column(2)	enum class time_frequency {}	struct KthValueVisitor { }	gen_negative_binomial_dist()
get_data_by_idx( )	struct BadRange { }	struct MaxVisitor { }	gen_normal_dist( )
get_data_by_loc( )	struct ColNotFound { }	struct MeanVisitor { }	gen_poisson_dist( )
get_data_by_rand( )	struct DataFrameError { }	struct MedianVisitor { }	gen_student_t_dist( )
get_data_by_sel(3)	struct InconsistentData { }	struct MinVisitor { }	gen_uniform_int_dist( )
get_index(2)	struct Index2D { }	struct ModeVisitor { }	gen_uniform_real_dist( )
get_row( )	struct NotFeasible { }	struct NLargestVisitor { }	gen_weibull_dist( )
get_view_by_idx( )	struct NotImplemented { }	struct NSmallestVisitor { }	
get_view_by_loc( )		struct ProdVisitor { }	
get_view_by_rand( )		struct ReturnVisitor { }	
get_view_by_sel(3)		struct SimpleRollAdopter { }	
groupby( )		struct SLRegressionVisitor { }	
groupby_async( )		struct StatsVisitor { }	
is_equal( )		struct StdVisitor { }	
join_by_index( )		struct SumVisitor { }	
load_column(3)		struct TrackingErrorVisitor { }	
load_data( )			
load_index(2)			
make_consistent( )			
modify_by_idx( )			
multi_visit( )			
read( )			
read_async( )			
remove_column( )			
remove_data_by_idx( )			
remove_data_by_loc( )			
remove_data_by_sel(3)			
rename_column( )			
replace(2)			
replace_async(2)			
replace_index( )			
rotate( )			
self_bucketize( )			
self_rotate( )			
self_shift( )			
shape( )			
shift( )			
shuffle( )			
single_act_visit(2)			
sort( )			
sort_async( )			
transpose( )			
value_counts( )			
visit(5)			
write( )			
write_async( )			

Motivation

Although Pandas has a spot-on interface and it is full of useful functionalities, it lacks performance and scalability. For example, it is hard to decipher high-frequency intraday data such as Options data or S&P500 constituents tick-by-tick data using Pandas. Another issue I have encountered often is the research is done using Python, because it has such tools as Pandas, but the execution in production is in C++ for its efficiency, reliability and scalability. Therefore, there is this translation, or sometimes a bridge, between research and executions.

Also, in this day and age, C++ needs a heterogeneous data container.

Mainly because of these factors, I implemented the C++ DataFrame.

This library could still have more functionalities compared with Pandas. I welcome all contributions from people with expertise, interest, and time to do it. I will add more functionalities from time to time, but currently my spare time is limited.

Views were recently added. This is a very interesting/useful concept that even Pandas doesn't have it currently. A view is a slice of a DataFrame that is a reference to the original DataFrame. It appears exactly the same as a DataFrame, but if you modify any data in the view, the original DataFrame will also be modified.

There are certain things you cannot do in views. For example, you cannot add to delete columns, extend the index column, …

For more understanding, look at this document further and/or the test files.

Visitors are the mechanism to run statistical algorithms. Most of DataFrame statistical algorithms are in "visitors". Visitor is the mechanism by which DataFrame passes data points to your algorithm. You can add your own algorithms to a visitor functor and extend DataFrame easily. There are two kinds of visitation mechanisms in DataFrame:

1) Regular visit (visit()). In this case DataFrame passes the given column(s) data points one-by-one to the visitor functor. This is convenient for algorithms that can operate per data point, such as correlation, variance etc.
2) Single-action visit (single_act_visit()). In this case a reference to the given column(s) are passed to the visitor functor at once. This is necessary for algorithms that need the whole data together, such as return, median, etc.

See this document and *dataframe_tester.cc* for more examples and documentation.

Random Generators I have added a series of convenient stand-alone functions to generate random numbers (it covers all C++ standard distributions). You can seamlessly use these routines to generate random DataFrame columns.

See this *document* and file *RandGen.h*

Code structure

The DataFrame library is almost a header-only library with a few small source files exceptions, *HeteroVector.cc and HeteroView.cc* and a few others. Also there is *DateTime.cc*.

Starting from the root directory;

include directory contains most of the code. It includes *.h* and *.tcc* files. The latter are C++ template code files (they are mostly located in the *Internals* subdirectory). The main header file is *DataFrame.h*. It contains the entire DataFrame object and its interface. There are comprehensive comments for each interface call in that file. The rest of the files there will show you how the sausage is made.
Include directory also contains subdirectories that contain mostly internal DataFrame implementation.
One exception, the *DateTime.h* is located in the *Utils* subdirectory

src directory contains Linux-only make files and a few subdirectories that contain various source codes.

test directory contains all the test source files, mocked data files, and test output files. The main test source file is *dataframe_tester.cc*. It contains test cases for all functionalities of DataFrame. It is not in a very organized structure. I plan to make the test cases more organized.

Build Instructions

Using plain make and make-files

Go to the root of the repository, where license file is, and execute *build_all.sh*. This will build the library and test executables for Linux flavors.

Using cmake

Please see README file. Thanks to @justinjk007, you should be able to build this in Linux, Windows, Mac, and more

EXAMPLE

This library is based on a heterogenous vector. The heterogeneity is achieved by using static STL or STL-like vectors. Since C++ is a strongly typed language, you still have to know your column types per container at compile time. You can add more columns with different types at any time to your container, but when analyzing the data at any given time you must know the column types.

This is an example of how to create a DataFrame, load data, and run an operation on it:

```cpp
using namespace hmdf;

// Defines a DataFrame with unsigned long index type that used std::vector
using MyDataFrame = StdDataFrame<unsigned long>;

MyDataFrame             df;
std::vector<int>        intvec = { 1, 2, 3, 4, 5 };
std::vector<double>     dblvec = { 1.2345, 2.2345, 3.2345, 4.2345, 5.2345 };
std::vector<double>     dblvec2 = { 0.998, 0.3456, 0.056, 0.15678, 0.00345,
                                    0.923, 0.06743, 0.1 };
std::vector<std::string> strvec = { "Some string", "some string 2", "some string 3",
                                    "some string 4", "some string 5" };
std::vector<unsigned long> ulgvec = { 1UL, 2UL, 3UL, 4UL, 5UL, 8UL, 7UL, 6UL }
std::vector<unsigned long> xulgvec = ulgvec;

// This is only one way of loading data into the DataFrame. There are
// many different ways of doing it. Please see DataFrame.h and dataframe_tester.cc
int rc = df.load_data(std::move(ulgvec),
                      std::make_pair("int_col", intvec),
                      std::make_pair("dbl_col", dblvec),
                      std::make_pair("dbl_col_2", dblvec2),
                      std::make_pair("str_col", strvec),
                      std::make_pair("ul_col", xulgvec));
// Sort the Frame by index
df.sort<MyDataFrame::IndexType, int, double, std::string>();
// Sort the Frame by column "dbl_col_2"
df.sort<double, int, double, std::string>("dbl_col_2");

// A functor to calculate mean, variance, skew, kurtosis, defined in
// DataFrameVisitors.h file
StatsVisitor<double>  stats_visitor;

// Calculate the stats on column "dbl_col"
df.visit<double>("dbl_col", stats_visitor);
```

View Example:

```cpp
std::vector<unsigned long>  idx =
    { 123450, 123451, 123452, 123450, 123455, 123450, 123449 };
std::vector<double>     d1 = { 1, 2, 3, 4, 5, 6, 7 };
std::vector<double>     d2 = { 8, 9, 10, 11, 12, 13, 14 };
std::vector<double>     d3 = { 15, 16, 17, 18, 19, 20, 21 };
std::vector<double>     d4 = { 22, 23, 24, 25 };
std::vector<std::string> s1 = { "11", "22", "33", "xx", "yy", "gg", "string" };
MyDataFrame             df;

df.load_data(std::move(idx),
            std::make_pair("col_1", d1),
            std::make_pair("col_2", d2),
            std::make_pair("col_3", d3),
            std::make_pair("col_4", d4),
            std::make_pair("col_str", s1));

using MyDataFrameView = DataFrameView<unsigned long>;
```

```
MyDataFrameView  dfv = df.get_view_by_loc<double, std::string>({ 3, 6 });

dfv.get_column<double>("col_3")[0] = 88.0;
std::cout << "After changing a value on view: "
          << dfv.get_column<double>("col_3")[0]
          << " == " << df.get_column<double>("col_3")[3]
          << std::endl;
```

For more code examples see file *dataframe_testr.cc*

Types

```
using size_type = typename std::vector<DataVec>::size_type;
```
size_type is the size type

```
using IndexType = I;
```
IndexType is the type of the index column

```
using IndexVecType = std::vector<I>;
```
IndexVecType is the type of the vector containing the index column

```
enum class nan_policy : bool  {
    pad_with_nans = true,
    dont_pad_with_nans = false
};
```
Enumerated type of Boolean type to specify whether data should be padded with NaN or not

```
enum class sort_state : bool  {
    sorted = true,
    not_sorted = false
};
```
Enumerated type of Boolean type to specify whether data is currently sorted or not

```
template<typename T>
struct Index2D  {
    T   begin {};
    T   end {};
};
```
It represents a range with begin and end within a continuous memory space

```
enum class shift_policy : unsigned char  {
    down = 1,    // Shift/rotate the content of all columns down,
                 // keep index unchanged
    up = 2,      // Shift/rotate the content of all columns up,
                 // keep index unchanged
};
```
This policy is relative to a tabular data structure
There is no right or left shift (like Pandas), because columns in DataFrame have no ordering. They can only be accessed by name

```
enum class fill_policy : unsigned char  {
    value = 1,
    fill_forward = 2,
```

```
    fill_backward = 3,
    linear_interpolate = 4,        // Using the index as X coordinate
    linear_extrapolate = 5         // Using the index as X coordinate
};
```
This policy determines how to fill missing values in the DataFrame
value: Fill all the missing values, in a given column, with the given value.
fill_forward: Fill the missing values, in a given column, with the last valid
                value before the missing value
fill_backward: Fill the missing values, in a given column, with the first valid
                value after the missing value
linear_interpolate:
linear_extrapolate:
Use the index column as X coordinate and the given column as Y coordinate
And do interpolation/extrapolation as follows:

$$Y = Y1 + \frac{X - X1}{X2 - X1} * (Y2 - Y1)$$

```
enum class drop_policy : unsigned char  {
    all = 1,                 // Remove row if all columns are nan
    any = 2,                 // Remove row if any column is nan
    threshold = 3            // Remove row if threshold number of columns are nan
};
```
This policy specifies what rows to drop/remove based on missing column data
all: Drop the row if all columns are missing
any: Drop the row if any column is missing
threshold: Drop the column if threshold number of columns are missing

```
enum class io_format : unsigned char  {
    csv = 1,
};
```
This specifies the I/O format for reading and writing to/from files, streams, etc.
Currently only CSV format is supported. The CSV format is as follows:
-- Any empty line or any line started with # will be ignored
-- A data line has the following format:
    <column name>:<number of data points>:<\<type\>>:data,data,…
    An example line would look like this:
    price:1001:<double>:23.456,24.56,…

```
enum class time_frequency : unsigned char  {
    annual = 1,
    monthly = 2,
    weekly = 3,
    daily = 4,
    hourly = 5,
    minutely = 6,
```

```cpp
    secondly = 7,
    millisecondly = 8,
    // microsecondly = 9,
    // nanosecondly = 10
};
```
This enum specifies time frequency for index generation and otherwise. The names are self-explanatory.

```cpp
enum class return_policy : unsigned char  {
    log = 1,
    percentage = 2,
    monetary = 3,
};
```
This policy specifies the type of return to be calculated
log: log(present / past)
percentage: (present - past) / past)
monetary: present - past

```cpp
enum class random_policy : unsigned char  {
    num_rows_with_seed = 1,  // Number of rows with specifying a seed
    num_rows_no_seed = 2,    // Number of rows with no seed specification
    frac_rows_with_seed = 3, // Fraction of rows with specifying a seed
    frac_rows_no_seed = 4,   // Fraction of rows with no seed specification
};
```
Specification for calling get_[data|view]_by_rand()
Number of rows means the n parameter is an positive integer specifying the number of rows to select
Fraction of rows means the n parameter is a positive real number [0:1] specifying a fraction of rows to select

---

```cpp
template<typename T>
struct  RandGenParams  {
    // Minimum value
    T   min_value { std::numeric_limits<T>::min() };
    // Maximum value
    T   max_value { std::numeric_limits<T>::max() };

    // Generator seed. By using the same seed, you can generate the
    // same random numbers each time. -1 means do not use a seed,
    // so each generation will be different.
    unsigned int        seed { (unsigned int) -1 };

    // The p distribution parameter (probability of generating true)
    double      prob_true { 0.5 };
    // The t or k distribution parameter (number of trials)
    std::size_t   t_dist { 1 };
```

```cpp
    // The μ distribution parameter (the mean of the distribution)
    double      mean { 1.0 };
    // the σ distribution parameter (standard deviation)
    double      std { 0 };
    // The λ distribution parameter (the rate parameter)
    double      lambda { 1.0 };
    // The α distribution parameter (shape, location)
    double      alpha { 1.0 };
    // The β distribution parameter (scale)
    double      beta { 1.0 };
    // The m distribution parameter (log-scale)
    double      m { 0 };
    // The s distribution parameter (shape)
    double      s { 1.0 };
    // The n distribution parameter (degrees of freedom)
    double      n { 1.0 };
    // degrees of freedom for fisher_f_distribution
    double      n2 { 1.0 };
};
```

This structure is used as a list of parameters for the stand-alone random number generators. Note, not all parameters are applicable to all generators. Please refer to each generator documentation in this document for list of applicable parameters.

```cpp
template<typename T, typename U>
struct type_declare;

template<typename U>
struct type_declare<HeteroVector, U>  { using type = std::vector<U>; };

template<typename U>
struct type_declare<HeteroView, U>  { using type = VectorView<U>; };
```

This is a spoofy way to declare a type at compile time dynamically. Here it is used in declaring a few different data structures depending whether we are a DataFrame or DataFrameView

---

```cpp
template<typename I, typename H>
class DataFrame;

template<typename I>
using StdDataFrame = DataFrame<I, HeteroVector>;

template<typename I>
using DataFrameView = DataFrame<I, HeteroView>;
```

DataFrame is a class that has; An index column of type I (timestamp, although it doesn't have to be time), and many other columns of different types. The storage used throughout is std::vector.

DataFrames could be instantiated in two different modes:

StdDataFrame is the standard fully functional data-frame.

DataFrameView is a referenced to a slice of another data-frame. Most of the functionalities of *StdDataFrame* is also available on the *DataFrameView*. But some functionalities such as adding/removing columns etc. are not allowable on views. If you change any of the data in a *DataFrameView* the corresponding data in the original *StdDataFrame* will also be changed.

Methods

In the following methods, "I" stands for the Index type and "H" stands for a Heterogenous vector type:

template\<typename T\>
std::vector\<T\> &create_column(const char *name);
It creates an empty column named "name"

T: Type of the column
Returns a reference to the vector for that column

---

void *remove_column(const char *name);*
It removes a column named name.
The actual data vector is not deleted, but the column is dropped from DataFrame

---

void rename_column (const char *from, const char *to);
It renames column named from to to. If column from does not exists, it throws an exception

---

template\<typename ... Ts\>
size_type &load_data(IndexVecType &&indices, Ts ... args);
This is the most generalized load function. It creates and loads an index and a variable number of columns. The index vector and all column vectors are "moved" to DataFrame.

Ts: The list of types for columns in args
indices: A vector of indices (timestamps) of type IndexType;
args: A variable list of arguments consisting of
std::pair(\<const char *name, std::vector\<T\> &&data\>).
Each pair represents a column data and its name
Returns number of items loaded

---

template\<typename ITR\>
size_type load_index(const ITR &begin, const ITR &end);
It copies the data from iterators begin to end into the index column

ITR: Type of the iterator
Returns number of items loaded

---

size_type load_index(*IndexVecType* &&idx);
It moves the idx vector into the index column.
Returns number of items loaded

---

static std::vector\<I\>
gen_datetime_index(const char *start_datetime,

```
                    const char *end_datetime,
                    time_frequency t_freq,
                    long increment = 1,
                    DT_TIME_ZONE tz = DT_TIME_ZONE::LOCAL);
```
This static method generates a date/time-based index vector that could be fed directly to one of the load methods. Depending on the specified frequency, it generates specific timestamps (see below).
It returns a vector of I timestamps.
Currently I could be any built-in numeric type or DateTime

start_datetime, end_datetime: They are the start/end date/times of
                                    requested timestamps.
    They must be in the following format:
    MM/DD/YYYY [HH[:MM[:SS[.MMM]]]]
t_freq: Specifies the timestamp frequency. Depending on the frequency,
        and I type specific timestamps are generated as follows:
    - I type of DateTime always generates timestamps of DateTime.
    - Annual, monthly, weekly, and daily frequencies generates YYYYMMDD
      timestamps.
    - Hourly, minutely, and secondly frequencies generates epoch
      timestamps (64 bit).
    - Millisecondly frequency generates nano-second since epoch
      timestamps (128 bit).
increment: Increment in the units of the frequency
tz: Time-zone of generated timestamps

 NOTE: It is the responsibility of the programmer to make sure I type is big enough to contain the frequency.

```
static std::vector<IndexType>
gen_sequence_index(const IndexType &start_value,
                    const IndexType &end_value,
                    long increment = 1);
```
This static method generates a vector of sequential values of IndexType that could be fed directly to one of the load methods.
The values are incremented by "increment".
The index type must be incrementable.
If by incrementing "start_value" by increment you would never reach "end_value", the behavior will be undefined.
It returns a vector of IndexType values.

start_value, end_value: Starting and ending values of IndexType.
                            Start value is included. End value is excluded.
increment: Increment by value

```
template<typename T, typename ITR>
```

size_type
load_column(const char *name,
            Index2D<const ITR &> range,
            nan_policy padding = nan_policy::pad_with_nans);
It copies the data from iterators begin to end to the named column. If column does
not exist, it will be created. If the column exist, it will be over written.

T: Type of data being copied
ITR: Type of the iterator
name: Name of the column
range: The begin and end iterators for data
padding: If true, it pads the data column with nan if it is shorter than the
        index column.
Returns number of items loaded

template<typename T>
size_type
load_column(const char *name,
            std::vector<T> &&data,
            nan_policy padding = nan_policy::pad_with_nans);

template<typename T>
size_type
load_column(const char *name,
            const std::vector<T> &data,
            nan_policy padding = nan_policy::pad_with_nans);
It moves or copies (depending on the version) the data to the named column in
DataFrame. If column does not exist, it will be created. If the column exist, it will
be over written.

T: Type of data being moved
name: Name of the column
padding: If true, it pads the data column with nan,
        if it is shorter than the index column.
Returns number of items loaded

size_type append_index(const IndexType &val);
It appends val to the end of the index column.
Returns number of items loaded

template<typename ITR>
size_type append_index(Index2D<const ITR &> range);
It appends the range begin to end to the end of the index column

ITR: Type of the iterator
range: The begin and end iterators for data

Returns number of items loaded

---

template<typename T>
size_type
append_column(const char *name,
                const T &val,
                nan_policy padding = nan_policy::pad_with_nans);
It appends val to the end of the named data column. If data column doesn't exist,
it throws an exception.

T: Type of the named data column
name: Name of the column
padding: If true, it pads the data column with nan,
        if it is shorter than the index column.
Returns number of items loaded

template<typename T, typename ITR>
size_type
append_column(const char *name,
                Index2D<const ITR &> range,
                nan_policy padding = nan_policy::pad_with_nans);
It appends the range begin to end to the end of the named data column. If data
column doesn't exist, it throws an exception.

T: Type of the named data column
ITR: Type of the iterator
name: Name of the column
range: The begin and end iterators for data
padding: If true, it pads the data column with nan,
        if it is shorter than the index column.
Returns number of items loaded

---

template<typename ... Ts>
void remove_data_by_idx (Index2D<I> range);
It removes the data rows from index begin to index end.
DataFrame must be sorted by index or behavior is undefined.
This function first calls make_consistent() that may add nan values to
data columns.

Ts: The list of types for all columns. A type should be specified only once.
range: The begin and end iterators for index specified with index values

---

template<typename ... Ts>
void remove_data_by_loc (Index2D<int> range);
It removes the data rows from location begin to location end
within range.

This function supports Python-like negative indexing. That is why the range type is int.
This function first calls make_consistent() that may add nan values to data columns.

Ts: The list of types for all columns. A type should be specified only once.
range: The begin and end iterators for data

```
template<typename T, typename F, typename ... Ts>
void remove_data_by_sel (const char *name, F &sel_functor);
```
It removes data rows by boolean filtering selection via the sel_functor (e.g. a functor, function, or lambda). Each element of the named column along with its corresponding index is passed to the sel_functor. If sel_functor returns true, that row will be removed.
The signature of sel_fucntor:
    bool ()(const IndexType &, const T &)

NOTE: If the selection logic results in empty column(s), the empty
        column(s) will _not_ be padded with NaN's. You can always call
        make_consistent() afterwards to make all columns into consistent length

T: Type of the named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name: Name of the data column
sel_functor: A reference to the selecting functor

```
template<typename T1, typename T2, typename F, typename ... Ts>
void
remove_data_by_sel (const char *name1, const char *name2, F &sel_functor);
```
This does the same function as above remove_data_by_sel() but operating on two columns.
The signature of sel_fucntor:
    bool ()(const IndexType &, const T1 &, const T2 &)

T1: Type of the first named column
T2: Type of the second named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name1: Name of the first data column
name2: Name of the second data column
sel_functor: A reference to the selecting functor

```
template<typename T1, typename T2, typename T3, typename F,
         typename ... Ts>
void
```

remove_data_by_sel (const char *name1, const char *name2, const char *name3,
F &sel_functor);
This does the same function as above remove_data_by_sel() but operating on
three columns.
The signature of sel_fucntor:
    bool ()(const IndexType &, const T1 &, const T2 &, const T3 &)

T1: Type of the first named column
T2: Type of the second named column
T3: Type of the third named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name1: Name of the first data column
name2: Name of the second data column
name3: Name of the third data column
sel_functor: A reference to the selecting functor

template<size_t N, typename ... Ts>
void
shuffle(const std::array<const char *, N> col_names, bool also_shuffle_index);
It randomly shuffles the named column(s) non-deterministically.

also_shuffle_index: If true, it shuffles the named column(s) and the
                    index column. Otherwise, index is not shuffled.
N: Number of named columns
Ts: The list of types for all columns. A type should be specified only once.

template<typename T, size_t N>
void fill_missing(const std::array<const char *, N> col_names,
                  fill_policy policy,
                  const std::array<T, N> values = { },
                  int limit = -1);
It fills all the "missing values" with the given values, and/or using the given
method (See fill_policy above). Missing is determined by being NaN for types
that have NaN. For types without NaN (e.g. string), default value is considered
missing value.

T: Type of the column(s) in col_names array
N: Size of col_names and values array
col_names: An array of names specifying the columns to fill.
policy: Specifies the method to use to fill the missing values.
        For example; forward fill, values, etc.
values: If the policy is "values", use these values to fill the missing
        holes. Each value corresponds to the same index in the
        col_names array.

limit: Specifies how many values to fill. Default is -1 meaning fill all missing values.

template<typename ... Ts>
void drop_missing(drop_policy policy, size_type threshold = 0);
It removes a row if any or all or some of the columns are NaN, based on drop policy

Ts: The list of types for all columns. A type should be specified only once.
threshold: If drop policy is threshold, it specifies the numbers of
            NaN columns before removing the row.

template<typename T, size_t N>
size_type replace(const char *col_name,
                const std::array<T, N> old_values,
                const std::array<T, N> new_values,
                int limit = -1);
It iterates over the column named col_name (or index, if col_name == "INDEX") and replaces all values in old_values with the corresponding values in new_values up to the limit. If limit is omitted, all values will be replaced.
It returns number of items replaced.

T: Type on column col_name. If this is index it would be the same as I.
N: Size of old_values and new_values arrays
col_name: Name of the column
old_array: An array of values to be replaced in col_name column
new_array: An array of values to replace the old_values in col_name column
limit: Limit of how many items to replace. Default is to replace all.

template<typename T, size_t N>
std::future<size_type>
replace_async(const char *col_name,
                const std::array<T, N> old_values,
                const std::array<T, N> new_values,
                int limit = -1);
Same as replace() above, but executed asynchronously
NOTE: multiple instances of replace_async() maybe executed for different columns at the same time with no problem.

template<typename T, typename F>
void replace(const char *col_name, F &functor);
This is similar to replace() above but it lets a functor replace the values in the named column. The functor is passed every value of the column along with a const reference of the corresponding index value.
Unlike the replace version above, this replace can only work on data columns. It will not work on index column.

The functor must have the following interface at minimum:
        bool operator() (const IndexType &ts, T &value);
A false return from the above operator method stops the iteration
through named column values.

T: Type on column col_name. If this is index it would be the same as I.
F: The functor type
col_name: Name of the column
functor: An instance of the functor

template<typename T, typename F>
std::future<void> replace_async(const char *col_name, F &functor);
Same as replace() above, but executed asynchronously
NOTE: multiple instances of replace_async() maybe executed for different
columns at the same time with no problem.

template<size_t N>
size_type
replace_index(const std::array<IndexType, N> old_values,
                const std::array<IndexType, N> new_values,
                int limit = -1);
This does the same thing as replace() above for the index column

N: Size of old_values and new_values arrays
old_array: An array of values to be replaced in col_name column
new_array: An array of values to replace the old_values in col_name column
limit: Limit of how many items to replace. Default is to replace all.

template<typename ... Ts>
void make_consistent ();
Make all data columns the same length as the index. If any data column is shorter
than the index column, it will be padded by nan.
This is also called by sort(), before sorting

Ts: The list of types for all columns. A type should be specified only once.

template<typename T, typename ... Ts>
void sort(const char *by_name = nullptr);
Sort the DataFrame by the named column. By default, it sorts by index (i.e.
by_name == nullptr). Sort first calls make_consistent() that may add nan values to
data columns. nan values make sorting nondeterministic.

T: Type of the by_name column. You always of the specify this type,
    even if it is being sorted to the default index
Ts: The list of types for all columns. A type should be specified only once.

```
template<typename T, typename ... Ts>
std::future<void> sort_async (const char *by_name = nullptr);
```
Same as sort() above, but executed asynchronously

```
template<typename F, typename T, typename ... Ts>
DataFrame
groupby (F &&func,
          const char *gb_col_name = nullptr,
          sort_state already_sorted = sort_state::not_sorted) const;
```
Groupby copies the DataFrame into a temp DataFrame and sorts the temp df by gb_col_name before performing groupby. If gb_col_name is null, it groups by index.

F: type functor to be applied to columns to group by
T: type of the groupby column. In case if index, it is type of index
    Ts: List of the types of all data columns.
    A type should be specified in the list only once.
Ts: The list of types for all columns. A type should be specified only once.
func: The functor to do the groupby. Specs for the functor is
        in a separate doc.
already_sorted: If the DataFrame is already sorted by gb_col_name,
                  this will save the expensive sort operation

```
std::future<DataFrame>
groupby_async (F &&func,
                const char *gb_col_name = nullptr,
                sort_state already_sorted = sort_state::not_sorted) const;
```
Same as groupby() above, but executed asynchronously

```
template<typename T>
StdDataFrame<T> value_counts (const char *col_name) const;
```
It counts the unique values in the named column.
It returns a StdDataFrame of following specs:
1) The index is of type T and contains all unique values in
    the named column.
2) There is only one column named "counts" of type size_type that
    contains the count for each index row.
For this method to compile and work, 3 conditions must be met:
1) Type T must be hashable. If this is a user defined type, you
    must enable and specialize std::hash.
 2) The equality operator (==) must be well defined for type T.
 3) Type T must match the actual type of the named column.
 Of course, if you never call this method in your application,
 you need not be worried about these conditions.

T: Type of the col_name column.

template<typename F, typename ... Ts>
DataFrame bucketize (F &&func, const IndexType &bucket_interval) const;
It bucketizes the data and index into bucket_interval's, based on index values and calls the functor for each bucket. The result of each bucket will be stored in a new DataFrame with same shape and returned. Every data bucket is guaranteed to be as wide as bucket_interval. This mean some data items at the end may not be included in the new bucketized DataFrame. The index of each bucket will be the last index in the original DataFrame that is less than bucket_interval away from the previous bucket
NOTE: The DataFrame must already be sorted by index.

F: type functor to be applied to columns to bucketize
Ts: The list of types for all columns. A type should be specified only once.
bucket_interval: Bucket interval is in the index's single value unit.
                 For example, if index is in minutes, bucket_interval
                 will be in the unit of minutes and so on.
already_sorted: If the DataFrame is already sorted by index,
                 this will save the expensive sort operation

template<typename F, typename ... Ts>
std::future<DataFrame>
bucketize_async (F &&func, const IndexType &bucket_interval) const;
Same as bucketize() above, but executed asynchronously

template<typename F, typename ... Ts>
void self_bucketize (F &&func, const IndexType &bucket_interval);
This is exactly the same as bucketize() above. The only difference is it stores the result in itself and returns void. So, after the return the original data is lost and replaced with bucketized data

template<typename T, typename V>
DataFrame
transpose(IndexVecType &&indices,
          const V &current_col_order,
          const V &new_col_names) const;
It transposes the data in the DataFrame.
The transpose() is only defined for DataFrame's that have a single data type.
NOTE: Since DataFrame columns have no ordering, the user must specify the order with current_col_order.

T: The single type for all data columns
V: The type of string vector specifying the new names for new columns
     after transpose
indices: A vector on indices for the new transposed DataFrame.

Its length must equal the number of rows in this DataFrame.
Otherwise an exception is thrown
current_col_order: A vector of strings specifying the order of columns
in the original DataFrame.
new_col_names: A vector of strings, specifying the column names for the
new transposed DataFrame.
Its length must equal the number of rows in this
DataFrame. Otherwise an exception is thrown

---

template<typename RHS_T, typename ... Ts>
StdDataFrame<I> join_by_index (const RHS_T &rhs, join_policy mp) const;
It joins the data between self (lhs) and rhs and returns the joined data
in a StdDataFrame, based on specification in join_policy.
The following conditions must be meet for this method
to compile and work properly:
1) I type must be the same between lhs and rhs.
2) Ordering (< > != ==) must be well defined for type I
3) Both lhs and rhs must be sorted by index
4) In both lhs and rhs, columns with the same name must have the same Type

RHS_T: Type of DataFrame rhs
Ts: The list of types for all columns. A type should be specified only once.
rhs: The rhs DataFrame
join_policy: Specifies how to join. For example inner join,
or left join, etc. (See join_policy definition)

---

template<typename ... Ts>
void self_shift (size_type periods, shift_policy sp);
It shifts all the columns in self up or down based on shift_policy.
Values that are shifted will be assigned to NaN. The index column remains
unchanged.
If user shifts with periods that is larger than the column length,
all values in that column become NaN.

Ts: The list of types for all columns. A type should be specified only once.
periods: Number of periods to shift
shift_policy: Specifies the direction (i.e. up/down) to shift

template<typename ... Ts>
StdDataFrame<I> shift (size_type periods, shift_policy sp) const;
It is exactly the same as self_shift, but it leaves self unchanged
and returns a new DataFrame with columns shifted.

---

template<typename ... Ts>
void self_rotate (size_type periods, shift_policy sp);
It rotates all the columns in self up or down based on shift_policy.

The index column remains unchanged.
If user rotates with periods that is larger than the column length, the behavior is undefined.

Ts: The list of types for all columns. A type should be specified only once.
periods: Number of periods to rotate
shift_policy: Specifies the direction (i.e. up/down) to rotate

template<typename ... Ts>
StdDataFrame<I> rotate (size_type periods, shift_policy sp) const;
It is exactly the same as self_rotate, but it leaves self unchanged
and returns a new DataFrame with columns rotated.

template<typename S, typename ... Ts>
bool write (S &o, bool values_only = false, io_format iof = io_format::csv) const;
It outputs the content of DataFrame into the stream o as text in the following format:
> *INDEX:<Comma delimited list of values>*
> *<Column1 name>:<Column1 type>:<Comma delimited list of values>*
> *<Column2 name>:<Column2 type>:<Comma delimited list of values>*

S: Output stream type
Ts: The list of types for all columns. A type should be specified only once.
o: Reference to an streamable object (e.g. cout)
values_only: If true, the name and type of each column is not written
iof: Specifies the I/O format. The default is CSV

template<typename S, typename ... Ts>
std::future<bool>
write_async (S &o,
             bool values_only = false,
             io_format iof = io_format::csv) const;
Same as write() above, but executed asynchronously

bool read (const char *file_name, io_format iof = io_format::csv);
It inputs the contents of a text file into itself (i.e. DataFrame). The format of the file must be:
> *INDEX:<Comma delimited list of values>*
> *<Column1 name>:<Column1 type>:<Comma delimited list of values>*
> *<Column2 name>:<Column2 type>:<Comma delimited list of values>*
All empty lines or lines starting with # will be skipped.

file_name: Complete path to the file
iof: Specifies the I/O format. The default is CSV

std::future<bool>

read_async (const char *file_name, io_format iof = io_format::csv);
Same as read() above, but executed asynchronously

---

std::pair<size_type, size_type> shape();
It returns a pair containing number of rows and columns.
Note: Number of rows is the number of index rows. Not every column
has the same number of rows, necessarily. But each column has,
at most, this number of rows.

---

template<typename T>
typename type_declare<H, T>::type &
get_column (const char *name);
It returns a reference to the container of named data column
The return type depends on if we are in standard or view mode

T: Data type of the named column

template<typename T>
const typename type_declare<H, T>::type &
get_column (const char *name) const;
It returns a const reference to the container of named data column
The return type depends on if we are in standard or view mode

T: Data type of the named column

---

template<size_t N, typename ... Ts>
HeteroVector
get_row(size_type row_num, const std::array<const char *, N> col_names) const;
It returns the data in row row_num for columns in col_names. The order of data
items in the returned vector is the same as order of columns on col_names.
The first item in the returned vector is always the index value
corresponding to the row_num
It returns a HeteroVector which contains a different type for each column.

N: Size of col_names and values array
Ts: The list of types for all columns. A type should be specified only once.
row_num: The row number
col_names: Names of columns to get data from. It also specifies the
order of data in the returned vector

---

template<typename T>
std::vector<T> get_col_unique_values (const char *name) const;
It returns a vector of unique values in the named column in the same
order that exists in the column.
For this method to compile and work, 3 conditions must be met:
1) Type T must be hash-able. If this is a user defined type, you

must enable and specialize std::hash.
2) The equality operator (==) must be well defined for type T.
3) Type T must match the actual type of the named column.
   Of course, if you never call this method in your application,
   you need not be worried about these conditions.

T: Data type of the named column

template<typename ... Ts>
DataFrame get_data_by_idx (Index2D<I> range) const;
It returns a DataFrame (including the index and data columns) containing the data from index begin to index end. This function assumes the DataFrame is consistent and sorted by index. The behavior is undefined otherwise.

Ts: The list of types for all columns. A type should be specified only once.
range: The begin and end iterators for index specified with index values

template<typename ... Ts>
DataFrameView<I> get_view_by_idx (Index2D<I> range) const;
It behaves like get_data_by_idx(), but it returns a DataFrameView.
A view is a DataFrame that is a reference to the original DataFrame.
So if you modify anything in the view the original DataFrame will
also be modified.
Note: There are certain operations that you cannot do with a view.
         For example, you cannot add/delete columns, etc.

Ts: The list of types for all columns. A type should be specified only once.
range: The begin and end iterators for index specified with index values

template<typename ... Ts>
DataFrame get_data_by_loc (Index2D<int> range) const;
It returns a DataFrame (including the index and data columns) containing the data from location begin to location end.
This function supports Python-like negative indexing. That is why the range type is int.
This function assumes the DataFrame is consistent and sorted by index. The behavior is undefined otherwise.

Ts: The list of types for all columns. A type should be specified only once.
range: The begin and end iterators for data

template<typename ... Ts>
DataFrameView<I> get_view_by_loc (Index2D<int> range) const;
It behaves like get_data_by_loc(), but it returns a DataFrameView.
A view is a DataFrame that is a reference to the original DataFrame.
So if you modify anything in the view the original DataFrame will

also be modified.
Note: There are certain operations that you cannot do with a view.
       For example, you cannot add/delete columns, etc.

Ts: The list of types for all columns. A type should be specified only once.
range: The begin and end iterators for data

template<typename T, typename F, typename ... Ts>
DataFrame get_data_by_sel (const char *name, F &sel_functor) const;
This method does Boolean filtering selection via the sel_functor (e.g. a functor,
function, or lambda). It returns a new DataFrame. Each element of the named
column along with its corresponding index is passed to the sel_functor. If
sel_functor returns true, that index is selected and all the elements of all column
for that index will be included in the returned DataFrame.
The signature of sel_fucntor:
    bool ()(const IndexType &, const T &)

NOTE: If the selection logic results in empty column(s), the result
         empty columns will _not_ be padded with NaN's. You can always
         call make_consistent() on the original or result DataFrame to make
         all columns into consistent length

T: Type of the named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name: Name of the data column
sel_functor: A reference to the selecting functor

template<typename T, typename F, typename ... Ts>
DataFramePtrView<IndexType>
get_view_by_sel (const char *name, F &sel_functor);
This is identical with above get_data_by_sel(), but:
1) The result is a view
2) Since the result is a view, you cannot call make_consistent() on the result.

T: Type of the named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name: Name of the data column
sel_functor: A reference to the selecting functor

template<typename T1, typename T2, typename F, typename ... Ts>
DataFrame
get_data_by_sel (const char *name1, const char *name2, F &sel_functor) const;
This does the same function as above get_data_be_sel() but operating on two
columns.

The signature of sel_fucntor:
 bool ()(const IndexType &, const T1 &, const T2 &)

T1: Type of the first named column
T2: Type of the second named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name1: Name of the first data column
name2: Name of the second data column
sel_functor: A reference to the selecting functor

template<typename T1, typename T2, typename F, typename ... Ts>
DataFramePtrView<IndexType>
get_view_by_sel (const char *name1, const char *name2, F &sel_functor);
This is identical with above get_data_by_sel(), but:
1) The result is a view
2) Since the result is a view, you cannot call make_consistent() on the result.

T1: Type of the first named column
T2: Type of the second named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name1: Name of the first data column
name2: Name of the second data column
sel_functor: A reference to the selecting functor

template<typename T1, typename T2, typename T3, typename F,
   typename ... Ts>
DataFrame
get_data_by_sel (const char *name1, const char *name2, const char *name3,
     F &sel_functor) const;
This does the same function as above get_data_be_sel() but operating on three
columns.
The signature of sel_fucntor:
 bool ()(const IndexType &, const T1 &, const T2 &, const T3 &)

T1: Type of the first named column
T2: Type of the second named column
T3: Type of the third named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name1: Name of the first data column
name2: Name of the second data column
name3: Name of the third data column
sel_functor: A reference to the selecting functor

template<typename T1, typename T2, typename T3, typename F,
        typename ... Ts>
DataFramePtrView<IndexType>
get_view_by_sel (const char *name1, const char *name2, const char *name3,
                F &sel_functor);
This is identical with above get_data_by_sel(), but:
1) The result is a view
2) Since the result is a view, you cannot call make_consistent() on the result.

T1: Type of the first named column
T2: Type of the second named column
T3: Type of the third named column
F: Type of the selecting functor
Ts: The list of types for all columns. A type should be specified only once
name1: Name of the first data column
name2: Name of the second data column
name3: Name of the third data column
sel_functor: A reference to the selecting functor

template<typename ... Ts>
DataFrame
get_data_by_rand (random_policy spec, double n, size_type seed = 0) const;
It returns a DataFrame (including the index and data columns) containing the data
from uniform random selection. random_policy determines the behavior of
method.
Note: The actual number of rows returned might be smaller than requested. That
is because the random process might produce the same number more than once.
Note: The columns in the result are not padded with NaN.

Ts: The list of types for all columns. A type should be specified only once.
random_policy: Please see random_policy in DataFrameTypes.h. It specifies how
                this function should proceed.
n: Depending on the random policy, it is either the number of rows to sample or a
    fraction of rows to sample. In case of fraction, for example 0.4 means 40% of
    rows.
seed: depending on the random policy, user could specify a seed. The same seed
        should always produce the same random selection.

template<typename ... Ts>
DataFramePtrView<IndexType>
get_view_by_rand (random_policy spec, double n, size_type seed = 0) const;
It behaves like get_data_by_rand(), but it returns a DataFrameView. A view is a
DataFrame that is a reference to the original DataFrame. So if you modify
anything in the view the original DataFrame will also be modified.
Note: There are certain operations that you cannot do with a view. For example,
you cannot add/delete columns, etc.

Note: The columns in the result are not padded with NaN.

Ts: The list of types for all columns. A type should be specified only once.
random_policy: Please see random_policy in DataFrameTypes.h. It specifies how this function should proceed.
n: Depending on the random policy, it is either the number of rows to sample or a fraction of rows to sample. In case of fraction, for example 0.4 means 40% of rows.
seed: depending on the random policy, user could specify a seed. The same seed should always produce the same random selection.

const IndexVecType &get_index () const { return (indices_); }
It returns a const reference to the index container

IndexVecType &get_index () { return (indices_); }
It returns a reference to the index container

template<typename ... Ts>
void multi_visit (Ts ... args);
This is the most generalized visit function. It visits multiple columns with the corresponding function objects sequentially. Each function object is passed every single value of the given column along with its name and the corresponding index value. All functions objects must have this signature
        *bool (const IndexType &i, const char *name, T &col_value)*
If the function object returns false, the DataFrame will stop iterating at that point on that column..
NOTE: This method could be used to implement a pivot table.

Ts: The list of types for columns in args
args: A variable list of arguments consisting of
        *std::pair(<const char *name,*
                *&std::function<bool (const IndexType &, const char *, T &)>)*.
        Each pair represents a column name and the functor to run on it.
        NOTE: The second member of pair is a _pointer_ to the function or
                functor object

template<typename T, typename V>
V &visit (const char *name, V &visitor);
It passes the values of each index and each named column to the functor visitor sequentially from beginning to end
NOTE: This method could be used to implement a pivot table.

T: Type of the named column
V: Type of the visitor functor
name: Name of the data column

template<typename T1, typename T2, typename V>
V &visit (const char *name1, const char *name2, V &visitor);
It passes the values of each index and the two named columns to the functor
visitor sequentially from beginning to end
NOTE: This method could be used to implement a pivot table.

T1: Type of the first named column
T2: Type of the second named column
V: Type of the visitor functor
name1: Name of the first data column
name2: Name of the second data column

template<typename T1, typename T2, typename T3, typename V>
V &visit (const char *name1, const char *name2, const char *name3, V &visitor);
It passes the values of each index and the three named columns to the functor
visitor sequentially from beginning to end
NOTE: This method could be used to implement a pivot table.

T1: Type of the first named column
T2: Type of the second named column
T3: Type of the third named column
V: Type of the visitor functor
name1: Name of the first data column
name2: Name of the second data column
name3: Name of the third data column

template<typename T1, typename T2, typename T3, typename T4, typename V>
V &visit (const char *name1,
          const char *name2,
          const char *name3,
          const char *name4,
          V &visitor);
It passes the values of each index and the four named columns to the functor
visitor sequentially from beginning to end
NOTE: This method could be used to implement a pivot table.

T1: Type of the first named column
T2: Type of the second named column
T3: Type of the third named column
T4: Type of the forth named column
V: Type of the visitor functor
name1: Name of the first data column
name2: Name of the second data column
name3: Name of the third data column
name4: Name of the fourth data column

```
template<typename T1, typename T2, typename T3, typename T4, typename T5,
        typename V>
V &visit (const char *name1,
        const char *name2,
        const char *name3,
        const char *name4,
        const char *name5,
        V &visitor);
```

It passes the values of each index and the five named columns to the functor visitor sequentially from beginning to end
NOTE: This method could be used to implement a pivot table.

T1: Type of the first named column
T2: Type of the second named column
T3: Type of the third named column
T4: Type of the fourth named column
T5: Type of the fifth named column
V: Type of the visitor functor
name1: Name of the first data column
name2: Name of the second data column
name3: Name of the third data column
name4: Name of the fourth data column
name5: Name of the fifth data column

```
template<typename T, typename V>
V &single_act_visit (const char *name, V &visitor);
```

This is similar to visit(), but it passes a const reference to the index vector and the named column vector at once the functor visitor. This is convenient for calculations that need the whole data vector, for example auto-correlation.

T: Type of the named column
V: Type of the visitor functor
name: Name of the data column

```
template<typename T1, typename T2, typename V>
V &single_act_visit (const char *name1, const char *name2, V &visitor);
```

This is similar to visit(), but it passes a const reference to the index vector and the two named column vectors at once the functor visitor. This is convenient for calculations that need the whole data vector.
NOTE: This method could be used to implement a pivot table.

T1: Type of the first named column
T2: Type of the second named column
V: Type of the visitor functor
name1: Name of the first data column
name2: Name of the second data column

```
template<typename ... Ts>
bool is_equal (const DataFrame &rhs) const;
```
It compares self with rhs. If both have the same indices, same number of columns, same names for each column, and all columns are equal, then it returns true. Otherwise it returns false

Ts: The list of types for all columns. A type should be specified only once.

```
template<typename ... Ts>
DataFrame &
modify_by_idx (DataFrame &rhs,
              sort_state already_sorted = sort_state::not_sorted);
```
It iterates over all indices in rhs and modifies all the data columns in self that correspond to the given index value. If not already_sorted, both rhs and self will be sorted by index. It returns a reference to self

Ts: The list of types for all columns. A type should be specified only once.
already_sorted: If the self and rhs are already sorted by index,
                 this will save the expensive sort operations

Global Operators

These are currently arithmetic operators declared in *include/DataFrame.h*. Because they all have to be templated, they cannot be defined as redefined built-in operators.

```
template<typename DF, typename ... Ts>
inline DF df_plus (const DF &lhs, const DF &rhs);

template<typename DF, typename ... Ts>
inline DF df_minus (const DF &lhs, const DF &rhs);

template<typename DF, typename ... Ts>
inline DF df_multiplies (const DF &lhs, const DF &rhs);

template<typename DF, typename ... Ts>
inline DF df_divides (const DF &lhs, const DF &rhs);
```
These arithmetic operations operate on the same-name and same-type columns on lhs and rhs. Each pair of entries is operated on, only if they have the same index value.
They return a new DataFrame
NOTE: Both lhs and rhs must be already sorted by index, otherwise the result is nonsensical.

Built-in Visitors

These are all defined in file *include/DataFrameVisitors.h*. Also see *test/data_frame_tester.cc* for example usage.

There are some common interfaces in most of the visitors. For example the following interfaces are common between most (but not all) visitors:

get_result()   -- It returns the result of the visitor/algo.

pre()   -- It is called by DataFrame each time before starting to pass the data to the visitor. pre() is the place to initialize the process

post()   -- It is called by DataFrame each time it is done with passing data to the visitor.

```
template<typename T,
         typename I = unsigned long,
         typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct MeanVisitor;
```

This functor class calculates the mean of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

```
template<typename T,
         typename I = unsigned long,
         typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct SumVisitor;
```

This functor class calculates the sum of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

```
template<typename T,
         typename I = unsigned long,
         typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct CumSumVisitor;
```

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class calculates the cumulative sum of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.
The result is a vector of running sums

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

```
template<typename T,
         typename I = unsigned long,
         typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct ProdVisitor;
```
This functor class calculates the product of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

```
template<typename T,
         typename I = unsigned long,
         typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct CumProdVisitor;
```
This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.
This functor class calculates the cumulative product of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.
See this document and datasci_tester.cc for examples.
The result is a vector of running products.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

```
template<typename T, typename I = unsigned long>
struct MaxVisitor;
```
This functor class calculates the maximum of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type

template<typename T, typename I = unsigned long>
struct CumMaxVisitor;

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class calculates the cumulative maximum of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.

The result is a vector of running maximums

T: Column/data type
I: Index type

---

template<typename T, typename I = unsigned long>
struct MinVisitor;

This functor class calculates the minimum of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type

---

template<typename T, typename I = unsigned long>
struct CumMinVisitor;

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class calculates the cumulative minimum of a given column. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.

The result is a vector of running minimum

T: Column/data type
I: Index type

---

template<std::size_t N, typename T, typename I = unsigned long>
struct NLargestVisitor;

This functor class calculates the N largest values of a column. I runs in $O(N*M)$, where N is the number of largest values and M is the total number of all values. If N is relatively small this better than $O(M*logM)$. The constructor takes a single optional Boolean argument to whether skip NaN values. The default is True.

See this document and datasci_tester.cc for examples.

N: Number of largest values
T: Column/data type

I: Index type

---

template<std::size_t N, typename T, typename I = unsigned long>
struct NSmallestVisitor;
This functor class calculates the N smallest values of a column. I runs in O(N*M),
where N is the number of largest values and M is the total number of all values.
If N is relatively small this is better than O(M*logM). The constructor takes a
single optional Boolean argument to whether skip NaN values. The default is
True.
See this document and datasci_tester.cc for examples.

N: Number of largest values
T: Column/data type
I: Index type

---

template<typename T,
        typename I = unsigned long,
        typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct CovVisitor;
This functor class calculates the covariance of two given columns. In addition, it
provides the variances of both columns.
        explicit CovVisitor (bool bias = true, bool skipnan = true);
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
        typename I = unsigned long,
        typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct VarVisitor;
This functor class calculates the variance of a given column.
        explicit VarVisitor (bool bias = true);
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
        typename I = unsigned long,
        typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct BetaVisitor;
This functor class calculates the beta (i.e. exposure) of the given first column to
the given second column (benchmark).

explicit BetaVisitor (bool bias = true);
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
       typename I = unsigned long,
       typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct StdVisitor;
This functor class calculates the standard deviation of a given column.
       explicit StdVisitor (bool bias = true);
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
       typename I = unsigned long,
       typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct TrackingErrorVisitor;
This functor class calculates the tracking error between two columns. Tracking error is the standard deviation of the difference vector.
       explicit TrackingErrorVisitor (bool bias = true);
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
       typename I = unsigned long,
       typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct CorrVisitor;
This functor class calculates the correlation of two given columns.
       explicit CorrVisitor (bool bias = true);
See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
       typename I = unsigned long,

typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct AutoCorrVisitor;

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class calculates the auto correlation of given column. The result is a vector of auto correlations with lags of 0 up to length of column – 4.

See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

template<typename T,
            typename I = unsigned long,
            typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct ReturnVisitor;

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class calculates the return of a given column, according to the return policy (monetary, percentage, or log). The result is a vector of returns.

  explicit ReturnVisitor (return_policy rp);

See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

template<typename T,
            typename I = unsigned long,
            typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct KthValueVisitor;

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class finds the Kth element in the given column in linear time.

  explicit KthValueVisitor (size_type ke, bool skipnan = true);

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

template<typename T,
            typename I = unsigned long,
            typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct MedianVisitor;

This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.

This functor class finds the median of the given column, using the above Kth element visitor. It computes in linear time.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

template<std::size_t N, typename T, typename I = unsigned long>
struct ModeVisitor;
This is a "single action visitor", meaning it is passed the whole data vector in one call and you must use the single_action_visit() interface.
This functor class finds the N highest mode (N most repeated values) of the given column.
The result is an array of N items each of this type:

```
struct  DataItem  {
    // Value of the column item
    value_type                        value { };
    // List of indices where value occurred
    std::vector<index_type>          indices { };
    // Number of times value occurred
    inline size_type repeat_count() const  { return (indices.size()); }
    // List of column indices where value occurred
    std::vector<size_type>          value_indices_in_col {  };
};
```

N: Number of modes to find
T: Column/data type
I: Index type
T must be an arithmetic-enabled type

template<typename T,
         typename I = unsigned long,
         typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct DotProdVisitor;
This functor class calculates the dot-product of two given columns. See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

template<typename F, typename T, typename I = unsigned long>
struct SimpleRollAdopter;
This functor applies functor *F* to the data in a rolling progression. The roll count is given to the constructor of *SimpleRollAdoptor*. The result is a vector of *F* results for roll count.

inline SimpleRollAdopter(F &&functor, size_t roll_count)

F: Functor type
T: Column/data type
I: Index type

---

template<typename T,
        typename I = unsigned long,
        typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct StatsVisitor;
This functor class calculates the following statistics of a given column; mean, variance, standard deviation, skew, and kurtosis. See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

---

template<typename T,
        typename I = unsigned long,
        typename = typename std::enable_if<std::is_arithmetic<T>::value, T>::type>
struct SLRRegressionVisitor;
This functor class calculates simple linear regression, in one pass, of two given columns (x, y). See this document and datasci_tester.cc for examples.

T: Column/data type
I: Index type
T must be an arithmetic-enabled type

RANDOM GENERATORS

A set of convenient routines to generate random number.

For the definition and defaults of *RandGenParams*, see this *document* and file *DataFrameTypes.h*

It generates n uniform integer distribution random numbers.

$$P(i|a,b) = \frac{1}{b - a + 1}$$

.

It returns the vector of results
Optional parameters to set:
   max_value, min_value, seed

T: Type can only be [unsigned] char, [unsigned] short, [unsigned] int,
   [unsigned] long int, or [unsigned] long long int
n: Number of numeric to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_uniform_int_dist(std::size_t n, const RandGenParams<T> &params = { });

It generates n uniform real distribution random numbers.

$$P(i|a,b) = \frac{1}{b - a}$$

.

It returns the vector of results
Optional parameters to set:
   max_value, min_value, seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_uniform_real_dist(std::size_t n, const RandGenParams<T> &params = { });

It generates random Boolean values, according to the discrete probability function. The probability of true is:

$$P(b|p) = \begin{cases} p \text{ if } b == true \\ 1 - p \text{ if } b == fasle \end{cases}$$

It returns the vector of results
Optional parameters to set:

prob_true, seed

n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
std::vector<bool>
gen_bernoulli_dist(std::size_t n, const RandGenParams<bool> &params = { });

Produces random non-negative integer values i, distributed according to discrete probability function:

$$P(i|t,p) = \binom{t}{i} \cdot p^i \cdot (1-p)^{t-i}$$

It returns the vector of results
Optional parameters to set:
   t_dist, prob_true, seed

T: Type can only be [unsigned] char, [unsigned] short, [unsigned] int,
   [unsigned] long int, or [unsigned] long long int
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_binomial_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random non-negative integer values i, distributed according to discrete probability function:

$$P(i|k,p) = \binom{k+i-1}{i} \cdot p^k \cdot (1-p)^i$$

It returns the vector of results
Optional parameters to set:
   t_dist, prob_true, seed

T: Type can only be [unsigned] char, [unsigned] short, [unsigned] int,
   [unsigned] long int, or [unsigned] long long int
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_negative_binomial_dist(std::size_t n, const RandGenParams<T> &params = {});

Produces random non-negative integer values i, distributed according to discrete probability function:

$$P(i|p) = p \cdot (1-p)^i$$

It returns the vector of results
Optional parameters to set:

prob_true, seed

T: Type can only be [unsigned] char, [unsigned] short, [unsigned] int,
    [unsigned] long int, or [unsigned] long long int
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_geometric_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random non-negative integer values i, distributed according to
discrete probability function:

$$P(i|u) = \frac{e^{-u} \, u^{i}}{i!}$$

It returns the vector of results
Optional parameters to set:
    mean, seed

T: Type can only be [unsigned] char, [unsigned] short, [unsigned] int,
    [unsigned] long int, or [unsigned] long long int
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_poisson_dist(std::size_t n, const RandGenParams<T> &params = { });

duces random non-negative floating-point values x,
distributed according to probability density function:
        $P(x|\lambda) = \lambda e - \lambda x$

It returns the vector of results
Optional parameters to set:
    lambda, seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_exponential_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random non-negative floating-point values x,
distributed according to probability density function:

$$p(x|\alpha,\beta) = e^{-x/\beta} \beta^{\alpha} \cdot \Gamma(\alpha) \cdot X^{\alpha-1}$$

It returns the vector of results
Optional parameters to set:
   alpha, beta, seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_gamma_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random non-negative floating-point values x,
distributed according to probability density function:
see *https://en.cppreference.com/w/cpp/numeric/random/weibull_distribution*

It returns the vector of results
Optional parameters to set:
   alpha (shape), beta (scale), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_weibull_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random numbers according to the extreme value distribution
(it is also known as Gumbel Type I, log-Weibull, Fisher-Tippett Type I):
see *https://en.cppreference.com/w/cpp/numeric/random/extreme_value_distribution*

It returns the vector of results
Optional parameters to set:
   alpha (location), beta (scale), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_extreme_value_dist(std::size_t n, const RandGenParams<T> &params = { });

Generates random numbers according to the Normal (or Gaussian)
random number distribution. It is defined as:
see *https://en.cppreference.com/w/cpp/numeric/random/normal_distribution*

It returns the vector of results
Optional parameters to set:
  mean, std (scale), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_normal_dist(std::size_t n, const RandGenParams<T> &params = { });

The lognormal_distribution random number distribution produces
random numbers x > 0 according to a log-normal distribution:
see *https://en.cppreference.com/w/cpp/numeric/random/lognormal_distribution*

It returns the vector of results
Optional parameters to set:
  m (log-scale), s (shape), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_lognormal_dist(std::size_t n, const RandGenParams<T> &params = { });

The chi_squared_distribution produces random numbers x > 0 according
to the Chi-squared distribution:
see *https://en.cppreference.com/w/cpp/numeric/random/chi_squared_distribution*

It returns the vector of results
Optional parameters to set:
  n (degree of freedom), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_chi_squared_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random numbers according to a Cauchy distribution
(also called Lorentz distribution):
see *https://en.cppreference.com/w/cpp/numeric/random/cauchy_distribution*

It returns the vector of results
Optional parameters to set:
   alpha (location), beta (scale), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_cauchy_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random numbers according to the f-distribution:
see *https://en.cppreference.com/w/cpp/numeric/random/fisher_f_distribution*

It returns the vector of results
Optional parameters to set:
   n (degree of freedom), n2 (degree of freedom), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_fisher_f_dist(std::size_t n, const RandGenParams<T> &params = { });

Produces random floating-point values x, distributed according
to probability density function:
see *https://en.cppreference.com/w/cpp/numeric/random/student_t_distribution*

It returns the vector of results
Optional parameters to set:
   n (degree of freedom), seed

T: Type can only be float, double, or long double
n: Number of numerics to generate
params: List of all applicable parameters, see DataFrameTypes.h
template<typename T>
std::vector<T>
gen_student_t_dist(std::size_t n, const RandGenParams<T> &params = { });